

# Chapter 5

## Data Transformation

Icaro Franco Hernandez

2022-09-26

Remember that to view a whole data set we can execute, for example, `view(nycflights13::flights)`. This is a tibble, and tibble  $\neq$  table. Tibbles work better for the tidyverse. To check what kind of variable we are working with, we can use the following command:

```
typeof(nycflights13::flights$time_hour)
```

```
## [1] "double"
```

## Filtering

Selecting all flights from January first:

```
nycflights13::flights%>%
```

```
  dplyr::filter(month==1,day==1)->jan1
```

*#remeber that dplyr does not change the original dataset (always try to be as pure as possible).*

If we want to also print the new data set, just put between parenthesis:

```
(nycflights13::flights%>%
```

```
  dplyr::filter(month==1,day==1)->jan1)
```

```
## # A tibble: 842 x 19
```

```
##   year month   day dep_time sched_de~1 dep_d~2 arr_t~3 sched~4 arr_d~5 carrier
##   <int> <int> <int>   <int>      <int>    <dbl>   <int>   <int>    <dbl> <chr>
## 1  2013     1     1     517        515         2     830     819        11 UA
## 2  2013     1     1     533        529         4     850     830        20 UA
## 3  2013     1     1     542        540         2     923     850        33 AA
## 4  2013     1     1     544        545        -1    1004    1022       -18 B6
## 5  2013     1     1     554        600        -6     812     837       -25 DL
## 6  2013     1     1     554        558        -4     740     728        12 UA
## 7  2013     1     1     555        600        -5     913     854        19 B6
## 8  2013     1     1     557        600        -3     709     723       -14 EV
## 9  2013     1     1     557        600        -3     838     846        -8 B6
## 10 2013     1     1     558        600        -2     753     745         8 AA
## # ... with 832 more rows, 9 more variables: flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1: sched_dep_time, 2: dep_delay, 3: arr_time, 4: sched_arr_time,
## #   5: arr_delay
```

Boolean operators are: `&` for **and**, `|` for **or** and `!` for **is not**.

Inclusion operator: `%in%`. For example:

```
nycflights13::flights%>%
  dplyr::filter(month %in% c(11,12))
```

```
## # A tibble: 55,403 x 19
##   year month   day dep_time sched_de~1 dep_d~2 arr_t~3 sched~4 arr_d~5 carrier
##   <int> <int> <int>   <int>   <int>   <dbl>   <int>   <int>   <dbl> <chr>
## 1  2013    11     1       5     2359       6     352     345       7 B6
## 2  2013    11     1      35     2250     105     123     2356      87 B6
## 3  2013    11     1     455       500      -5     641     651     -10 US
## 4  2013    11     1     539       545      -6     856     827      29 UA
## 5  2013    11     1     542       545      -3     831     855     -24 AA
## 6  2013    11     1     549       600     -11     912     923     -11 UA
## 7  2013    11     1     550       600     -10     705     659       6 US
## 8  2013    11     1     554       600      -6     659     701      -2 US
## 9  2013    11     1     554       600      -6     826     827      -1 DL
## 10 2013    11     1     554       600      -6     749     751      -2 DL
## # ... with 55,393 more rows, 9 more variables: flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1: sched_dep_time, 2: dep_delay, 3: arr_time, 4: sched_arr_time,
## #   5: arr_delay
```

This will filter all flights that happened in november **or** december. The filter already excludes NA values.

## Exercises 5.2.4

1

Find all flights that

- Had an arrival delay of two or more hours

```
nycflights13::flights%>%
  dplyr::filter(arr_delay>=120)->f1
```

- Flew to Houston (IAH or HOU)

```
nycflights13::flights%>%
  dplyr::filter(dest %in% c("IAH", "HOU"))->f2
```

- Were operated by United, American, or Delta

```
nycflights13::flights%>%
  dplyr::filter(carrier %in% c("UA", "AA", "DL"))->f3
```

- Departed in Summer (July, August, and September)

```
nycflights13::flights%>%
  dplyr::filter(month %in% c(7, 8, 9))->fsummer
```

- Arrived more than two hours late, but did not leave late

```
nycflights13::flights%>%
  dplyr::filter(arr_delay>120, dep_time<=sched_dep_time)->f5
```

- Were delayed by at least an hour, but made up over 30 minutes in flight

```
nycflights13::flights%>%
  dplyr::filter(arr_delay>=60, air_time>30)->f6
```

- Departed between midnight and 6am (inclusive)

```
nycflights13::flights%>%
  dplyr::filter(hour %in% c(seq(0,6)))->f7
```

or

```
## function (e1, e2) .Primitive("|")
```

```
nycflights13::flights%>%
  dplyr::filter(hour >= 0 & hour<= 6)->f71
```

```
# nycflights13::flights%>%
#   dplyr::filter(0 <= hour <= 6) -> this does not work!
```

## 2

Another useful `dplyr` filter helper is `between()`. What does it do? Can you use it to simplify the code needed to answer the precious questions?

According to the documentation, `between()` let us pick any values between to boundaries, and it is a shortcut for `x>= & x<=`. They would be useful in the cases where we had to filter for the summer months and the flights between midnight and 6 a.m.:

```
nycflights13::flights%>%
  dplyr::filter(between(month, 7, 9))->f8
```

```
nycflights13::flights%>%
  dplyr::filter(between(hour, 0, 6))->f9
```

## 3

How many flights have a missing `dep_time`? What other variables are missing? What might these rows represent?

```
nycflights13::flights%>%
  dplyr::filter(is.na(dep_time))%>%
  dplyr::summarise(n = dplyr::n())->na
```

na

```
## # A tibble: 1 x 1
##       n
##   <int>
## 1  8255
```

Using the count operator from `dplyr` we can see that 8255 flights have missing values for the departure time. This means that these flights were canceled. If we do not have the departure time, we also cannot check the airtime, the departure delay and the arrival delay. **Remember this count operator (within the summarise function) from dplyr.**

## 4

Why is `NA^0` not missing? Why is `NA|TRUE` not missing? Why is `FALSE & NA` not missing? Can you figure out the general rule? (`NA*0` is a tricky counterexample!)

```
NA^0
```

```
## [1] 1
```

```
NA|TRUE
```

```
## [1] TRUE
```

```
FALSE&NA
```

```
## [1] FALSE
```

Since we are working with boolean operators here, the general rule is that R avoids the NA values and does let them contaminate the operation. It is different from the case if we calculate the average of some values with an NA (in that case it does contaminate the average).

```
v1<-c(1,1, NA)
mean(v1)
```

```
## [1] NA
```

```
mean(v1, na.rm = T)
```

```
## [1] 1
```

The command `na.rm=TRUE` discards the NA values from the calculation!

## Arranging

### Exercises 5.3.1

1

How could you use `arrange()` to sort all missing values to the start? (Hint: use `is.na()`)

```
v<-tibble::tibble(
  x=c(3,7,1,NA))
)
```

```
x=c(3,7,1,NA)
```

```
sort(x, decreasing = FALSE, na.last=FALSE)
```

```
## [1] NA 1 3 7
```

I used the base R command for sorting.

2

Sort `flights` to find the most delayed flights. Find the flights that left earliest.

For the most delayed flights we just need to arrange in descending format:

```
nycflights13::flights%>%
  dplyr::arrange(desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_de~1 dep_d~2 arr_t~3 sched~4 arr_d~5 carrier
##   <int> <int> <int>   <int>   <int>   <dbl>   <int>   <int>   <dbl> <chr>
## 1  2013     1     9     641     900    1301    1242    1530    1272 HA
## 2  2013     6    15    1432    1935    1137    1607    2120    1127 MQ
## 3  2013     1    10    1121    1635    1126    1239    1810    1109 MQ
## 4  2013     9    20    1139    1845    1014    1457    2210    1007 AA
## 5  2013     7    22     845    1600    1005    1044    1815     989 MQ
```

```
## 6 2013 4 10 1100 1900 960 1342 2211 931 DL
## 7 2013 3 17 2321 810 911 135 1020 915 DL
## 8 2013 6 27 959 1900 899 1236 2226 850 DL
## 9 2013 7 22 2257 759 898 121 1026 895 DL
## 10 2013 12 5 756 1700 896 1058 2020 878 AA
## # ... with 336,766 more rows, 9 more variables: flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1: sched_dep_time, 2: dep_delay, 3: arr_time, 4: sched_arr_time,
## #   5: arr_delay
```

And for the earliest departures we just need to arrange them:

```
nycflights13::flights%>%
  dplyr::arrange(dep_delay)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_de-1 dep_d-2 arr_t-3 sched-4 arr_d-5 carrier
##   <int> <int> <int>   <int>   <int>   <dbl>   <int>   <int>   <dbl> <chr>
## 1  2013    12     7    2040     2123    -43     40    2352     48 B6
## 2  2013     2     3    2022     2055    -33    2240    2338    -58 DL
## 3  2013    11    10    1408     1440    -32    1549    1559    -10 EV
## 4  2013     1    11    1900     1930    -30    2233    2243    -10 DL
## 5  2013     1    29    1703     1730    -27    1947    1957    -10 F9
## 6  2013     8     9     729     755    -26    1002     955     7 MQ
## 7  2013    10    23    1907     1932    -25    2143    2143     0 EV
## 8  2013     3    30    2030     2055    -25    2213    2250    -37 MQ
## 9  2013     3     2    1431     1455    -24    1601    1631    -30 9E
## 10 2013     5     5     934     958    -24    1225    1309    -44 B6
## # ... with 336,766 more rows, 9 more variables: flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1: sched_dep_time, 2: dep_delay, 3: arr_time, 4: sched_arr_time,
## #   5: arr_delay
```

### 3

Sort flights to find the fastest (highest speed) flights

```
nycflights13::flights%>%
  dplyr::select(air_time,
                dplyr::everything())%>%
  dplyr::arrange(air_time)
```

```
## # A tibble: 336,776 x 19
##   air_time year month   day dep_time sched_d-1 dep_d-2 arr_t-3 sched-4 arr_d-5
##   <dbl> <int> <int> <int>   <int>   <int>   <dbl>   <int>   <int>   <dbl>
## 1      20  2013     1    16    1355     1315     40    1442    1411     31
## 2      20  2013     4    13     537     527     10     622     628     -6
## 3      21  2013    12     6     922     851     31    1021     954     27
## 4      21  2013     2     3    2153    2129     24    2247    2224     23
## 5      21  2013     2     5    1303    1315    -12    1342    1411    -29
## 6      21  2013     2    12    2123    2130     -7    2211    2225    -14
## 7      21  2013     3     2    1450    1500    -10    1547    1608    -21
## 8      21  2013     3     8    2026    1935     51    2131    2056     35
## 9      21  2013     3    18    1456    1329     87    1533    1426     67
```

```
## 10      21 2013      3      19      2226      2145      41      2305      2246      19
## # ... with 336,766 more rows, 9 more variables: carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1: sched_dep_time, 2: dep_delay, 3: arr_time, 4: sched_arr_time,
## #   5: arr_delay
```

4

Which flights traveled the farthest? Which traveled the shortest?

The ones traveled the farthest:

```
nycflights13::flights%>%
  dplyr::select(distance,
                dplyr::everything())%>%
  dplyr::arrange(desc(distance))
```

```
## # A tibble: 336,776 x 19
##   distance year month   day dep_time sched_d-1 dep_d-2 arr_t-3 sched-4 arr_d-5
##   <dbl> <int> <int> <int>   <int>   <int>   <dbl>   <int>   <int>   <dbl>
## 1    4983  2013     1     1     857     900     -3    1516    1530    -14
## 2    4983  2013     1     2     909     900      9    1525    1530     -5
## 3    4983  2013     1     3     914     900     14    1504    1530    -26
## 4    4983  2013     1     4     900     900      0    1516    1530    -14
## 5    4983  2013     1     5     858     900     -2    1519    1530    -11
## 6    4983  2013     1     6    1019     900     79    1558    1530     28
## 7    4983  2013     1     7    1042     900    102    1620    1530     50
## 8    4983  2013     1     8     901     900      1    1504    1530    -26
## 9    4983  2013     1     9     641     900   1301    1242    1530   1272
## 10   4983  2013     1    10     859     900     -1   1449    1530    -41
## # ... with 336,766 more rows, 9 more variables: carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1: sched_dep_time, 2: dep_delay, 3: arr_time, 4: sched_arr_time,
## #   5: arr_delay
```

The ones that traveled the shortest:

```
nycflights13::flights%>%
  dplyr::select(distance,
                dplyr::everything())%>%
  dplyr::arrange(distance)
```

```
## # A tibble: 336,776 x 19
##   distance year month   day dep_time sched_d-1 dep_d-2 arr_t-3 sched-4 arr_d-5
##   <dbl> <int> <int> <int>   <int>   <int>   <dbl>   <int>   <int>   <dbl>
## 1      17  2013     7    27      NA     106      NA      NA     245      NA
## 2      80  2013     1     3    2127    2129     -2    2222    2224     -2
## 3      80  2013     1     4    1240    1200     40    1333    1306     27
## 4      80  2013     1     4    1829    1615    134    1937    1721    136
## 5      80  2013     1     4    2128    2129     -1    2218    2224     -6
## 6      80  2013     1     5    1155    1200     -5    1241    1306    -25
## 7      80  2013     1     6    2125    2129     -4    2224    2224      0
## 8      80  2013     1     7    2124    2129     -5    2212    2224    -12
## 9      80  2013     1     8    2127    2130     -3    2304    2225     39
## 10     80  2013     1     9    2126    2129     -3    2217    2224     -7
```

```
## # ... with 336,766 more rows, 9 more variables: carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1: sched_dep_time, 2: dep_delay, 3: arr_time, 4: sched_arr_time,
## #   5: arr_delay
```

## Selecting Columns

Selecting by the specific name of each column:

```
nycflights13::flights%>%
  dplyr::select(
    year,
    month,
    day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

Selecting an interval of columns:

```
nycflights13::flights%>%
  dplyr::select(
    month:arr_time
  )
```

```
## # A tibble: 336,776 x 6
##   month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int>   <int>         <int>         <dbl>   <int>
## 1     1     1     517           515           2     830
## 2     1     1     533           529           4     850
## 3     1     1     542           540           2     923
## 4     1     1     544           545          -1    1004
## 5     1     1     554           600          -6     812
## 6     1     1     554           558          -4     740
## 7     1     1     555           600          -5     913
## 8     1     1     557           600          -3     709
## 9     1     1     557           600          -3     838
## 10    1     1     558           600          -2     753
## # ... with 336,766 more rows
```

Negative selection:

```
nycflights13::flights%>%
  dplyr::select(
    -(month:arr_time))
```

```
## # A tibble: 336,776 x 13
##   year sched_arr~1 arr_d~2 carrier flight tailnum origin dest air_t~3 dista~4
##   <int>      <int>   <dbl> <chr>   <int> <chr>   <chr> <chr>   <dbl>   <dbl>
## 1 2013         819     11 UA      1545 N14228 EWR    IAH     227    1400
## 2 2013         830     20 UA      1714 N24211 LGA    IAH     227    1416
## 3 2013         850     33 AA      1141 N619AA JFK    MIA     160    1089
## 4 2013        1022    -18 B6       725 N804JB JFK    BQN     183    1576
## 5 2013         837    -25 DL       461 N668DN LGA    ATL     116     762
## 6 2013         728     12 UA      1696 N39463 EWR    ORD     150     719
## 7 2013         854     19 B6       507 N516JB EWR    FLL     158    1065
## 8 2013         723    -14 EV      5708 N829AS LGA    IAD      53     229
## 9 2013         846     -8 B6        79 N593JB JFK    MCO     140     944
## 10 2013         745      8 AA       301 N3ALAA LGA    ORD     138     733
## # ... with 336,766 more rows, 3 more variables: hour <dbl>, minute <dbl>,
## #   time_hour <dtm>, and abbreviated variable names 1: sched_arr_time,
## #   2: arr_delay, 3: air_time, 4: distance
```

The helpers are from tidyselect:

- tidyselect::starts\_with("")
- tidyselect::ends\_with("")
- tidyselect::contains("")
- tidyselect::matches("")
- tidyselect::num\_range("")

We can select a few columns and also the rest of them:

```
nycflights13::flights%>%
  dplyr::select(
    minute,
    distance,
    dplyr::everything())
```

```
## # A tibble: 336,776 x 19
##   minute distance year month day dep_time sched_de~1 dep_d~2 arr_t~3 sched~4
##   <dbl>   <dbl> <int> <int> <int>   <int>      <int>   <dbl>   <int>   <int>
## 1    15    1400  2013     1     1     517        515     2     830     819
## 2    29    1416  2013     1     1     533        529     4     850     830
## 3    40    1089  2013     1     1     542        540     2     923     850
## 4    45    1576  2013     1     1     544        545    -1    1004    1022
## 5     0     762  2013     1     1     554        600    -6     812     837
## 6    58     719  2013     1     1     554        558    -4     740     728
## 7     0    1065  2013     1     1     555        600    -5     913     854
## 8     0     229  2013     1     1     557        600    -3     709     723
## 9     0     944  2013     1     1     557        600    -3     838     846
## 10    0     733  2013     1     1     558        600    -2     753     745
## # ... with 336,766 more rows, 9 more variables: arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   hour <dbl>, time_hour <dtm>, and abbreviated variable names
## #   1: sched_dep_time, 2: dep_delay, 3: arr_time, 4: sched_arr_time
```



## Exercises 5.4.1

1

Brainstorm as many ways as possible to select `dep_time`, `dep_delay`, `arr_time`, and `arr_delay` from `flights`.

2

What happens if you include the name of a variable multiple times in a `select()` call?

```
nycflights13::flights%>%
  dplyr::select(
    year,
    minute,
    year) #crime ocorre e nada acontece nesse caso
```

```
## # A tibble: 336,776 x 2
##   year minute
##   <int> <dbl>
## 1  2013     15
## 2  2013     29
## 3  2013     40
## 4  2013     45
## 5  2013      0
## 6  2013     58
## 7  2013      0
## 8  2013      0
## 9  2013      0
## 10 2013      0
## # ... with 336,766 more rows
```

```
nycflights13::flights%>%
  dplyr::select(
    year,
    -(year)) #empty vector
```

```
## # A tibble: 336,776 x 0
```

```
nycflights13::flights%>%
  dplyr::select(
    (year:day),
    year)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
```

```
## # ... with 336,766 more rows
```

If we include the same column more than one time, the selection takes into account just the first operation. Unless it is a contradictory selection, which is this case we probably get an empty vector.

### 3

What does the `any_of()` function do? Why might it be helpful in conjunction with this vector?

Lets read the documentation:

```
?any_of()
```

```
## starting httpd help server ... done
```

```
vars<-c("year","month","day","dep_delay","arr_delay")
```

```
nycflights13::flights%>%  
  dplyr::select(  
    any_of(vars))
```

```
## # A tibble: 336,776 x 5
```

```
##   year month   day dep_delay arr_delay  
##   <int> <int> <int>     <dbl>     <dbl>  
## 1  2013     1     1         2         11  
## 2  2013     1     1         4         20  
## 3  2013     1     1         2         33  
## 4  2013     1     1        -1        -18  
## 5  2013     1     1        -6        -25  
## 6  2013     1     1        -4         12  
## 7  2013     1     1        -5         19  
## 8  2013     1     1        -3        -14  
## 9  2013     1     1        -3         -8  
## 10 2013     1     1        -2          8
```

```
## # ... with 336,766 more rows
```

```
nycflights13::flights%>%  
  dplyr::select(  
    -any_of(vars))
```

```
## # A tibble: 336,776 x 14
```

```
##   dep_time sched_~1 arr_t~2 sched~3 carrier flight tailnum origin dest  air_t~4  
##   <int>    <int>    <int>    <int> <chr>    <int> <chr>    <chr> <chr>    <dbl>  
## 1     517      515      830      819 UA      1545 N14228 EWR   IAH      227  
## 2     533      529      850      830 UA      1714 N24211 LGA   IAH      227  
## 3     542      540      923      850 AA      1141 N619AA  JFK   MIA      160  
## 4     544      545     1004     1022 B6      725 N804JB  JFK   BQN      183  
## 5     554      600      812      837 DL      461 N668DN  LGA   ATL      116  
## 6     554      558      740      728 UA      1696 N39463 EWR   ORD      150  
## 7     555      600      913      854 B6      507 N516JB  EWR   FLL      158  
## 8     557      600      709      723 EV      5708 N829AS  LGA   IAD       53  
## 9     557      600      838      846 B6       79 N593JB  JFK   MCO      140  
## 10     558      600      753      745 AA      301 N3ALAA  LGA   ORD      138
```

```
## # ... with 336,766 more rows, 4 more variables: distance <dbl>, hour <dbl>,  
## #   minute <dbl>, time_hour <dtm>, and abbreviated variable names  
## #   1: sched_dep_time, 2: arr_time, 3: sched_arr_time, 4: air_time
```

The `any_of()` command takes a vector of variables and selects columns according to the variables inside the

vector. The documentation states that it is useful for negative selections and also that it does not check for errors. It is not clear yet to me what are the advantages.

4

Does the result of running the following code surprise you? How do the select helpers deal with case by default? How can you change that default?

```
nycflights13::flights%>%
  dplyr::select(contains("TIME"))
```

```
## # A tibble: 336,776 x 6
##   dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
##   <int>         <int>    <int>         <int>      <dbl> <dtm>
## 1      517           515      830           819      227 2013-01-01 05:00:00
## 2      533           529      850           830      227 2013-01-01 05:00:00
## 3      542           540      923           850      160 2013-01-01 05:00:00
## 4      544           545     1004          1022      183 2013-01-01 05:00:00
## 5      554           600      812           837      116 2013-01-01 06:00:00
## 6      554           558      740           728      150 2013-01-01 05:00:00
## 7      555           600      913           854      158 2013-01-01 06:00:00
## 8      557           600      709           723       53 2013-01-01 06:00:00
## 9      557           600      838           846      140 2013-01-01 06:00:00
## 10     558           600      753           745      138 2013-01-01 06:00:00
## # ... with 336,766 more rows
```

```
select(flights,
  tidysselect::contains("TIME", ignore.case=FALSE))
```

```
## # A tibble: 336,776 x 0
```

```
select(flights,
  tidysselect::contains("time", ignore.case=FALSE)) #this one should return the same results as the
```

```
## # A tibble: 336,776 x 6
##   dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
##   <int>         <int>    <int>         <int>      <dbl> <dtm>
## 1      517           515      830           819      227 2013-01-01 05:00:00
## 2      533           529      850           830      227 2013-01-01 05:00:00
## 3      542           540      923           850      160 2013-01-01 05:00:00
## 4      544           545     1004          1022      183 2013-01-01 05:00:00
## 5      554           600      812           837      116 2013-01-01 06:00:00
## 6      554           558      740           728      150 2013-01-01 05:00:00
## 7      555           600      913           854      158 2013-01-01 06:00:00
## 8      557           600      709           723       53 2013-01-01 06:00:00
## 9      557           600      838           846      140 2013-01-01 06:00:00
## 10     558           600      753           745      138 2013-01-01 06:00:00
## # ... with 336,766 more rows
```

The result does surprise. At a first glance we would thought that this selection would return empty vectors, however the default setting of `contains()` ignore the differences between upper and lower case. If we want to make a more precise selection (and make a distinction of upper and lower case) we can add `ignore.case=FALSE`.

## Mutate

```
nycflights13::flights%>%
  dplyr::select(
    year:day,
    tidyrselect::ends_with("delay"),
    distance,
    air_time)%>%
  dplyr::mutate(
    gain=dep_delay - arr_delay,
    speed=(distance/air_time)*60)->flights_sml
```

If we want just to keep the new variables we can use `dplyr::transmute` instead.

We can use `dplyr::mutate` with any vectorised operation. If one vector is shorter than the other, the operation will be recycled (will repeat until the the end of the bigger vector).

- Integer division: `%/%`;
- Remainder: `%%`;
- Lead and lag: `lead()`, `lag()`;
- Cumulative mean: `cummean()`;

### Exercises 5.5.2

1

Currently `dep_time()` and `sched_dep_time` are convenient to look at, but hard to compute because they're not really continuous numbers. Convert them to a more convenient representation of number of minutes since midnight.

```
nycflights13::flights%>%
  dplyr::select(
    dep_time,
    sched_dep_time)%>%
  dplyr::mutate(
    dep_time_min=(dep_time%/%100)*60+(dep_time)%%100,
    sched_dep_time_min=(sched_dep_time%/%100)*60+(sched_dep_time)%%100)
```

```
## # A tibble: 336,776 x 4
##   dep_time sched_dep_time dep_time_min sched_dep_time_min
##   <int>      <int>      <dbl>      <dbl>
## 1      517        515        317        315
## 2      533        529        333        329
## 3      542        540        342        340
## 4      544        545        344        345
## 5      554        600        354        360
## 6      554        558        354        358
## 7      555        600        355        360
## 8      557        600        357        360
## 9      557        600        357        360
## 10     558        600        358        360
## # ... with 336,766 more rows
```

2

Compare `air_time` with `arr_time-dep_time`. What do you expect to see? What do you see? What do you need to do to fix it?

```
nycflights13::flights%>%
  dplyr::select(
    air_time,
    arr_time,
    dep_time)%>%
  dplyr::mutate(
    diff=arr_time-dep_time)
```

```
## # A tibble: 336,776 x 4
##   air_time arr_time dep_time diff
##   <dbl>    <int>    <int> <int>
## 1     227      830      517   313
## 2     227      850      533   317
## 3     160      923      542   381
## 4     183     1004      544   460
## 5     116      812      554   258
## 6     150      740      554   186
## 7     158      913      555   358
## 8      53      709      557   152
## 9     140      838      557   281
## 10    138      753      558   195
## # ... with 336,766 more rows
```

The variable air time is the amount of time spent in the air, in minutes. However the arrival and departure time are in the format of HoursMinutes. So, to fix it, we must compute the difference in minutes (or convert the arrival and departure time into minutes before subtracting).

```
nycflights13::flights%>%
  dplyr::select(
    air_time,
    arr_time,
    dep_time)%>%
  dplyr::mutate(
    arr_time_min = (arr_time%%100)*60+(arr_time%100),
    dep_time_min = (dep_time%%100)*60+(dep_time%100),
    air_time2 = arr_time_min-dep_time_min)
```

```
## # A tibble: 336,776 x 6
##   air_time arr_time dep_time arr_time_min dep_time_min air_time2
##   <dbl>    <int>    <int>    <dbl>        <dbl>    <dbl>
## 1     227      830      517        510          317        193
## 2     227      850      533        530          333        197
## 3     160      923      542        563          342        221
## 4     183     1004      544        604          344        260
## 5     116      812      554        492          354        138
## 6     150      740      554        460          354        106
## 7     158      913      555        553          355        198
## 8      53      709      557        429          357         72
## 9     140      838      557        518          357        161
## 10    138      753      558        473          358        115
## # ... with 336,766 more rows
```

### 3

Compare `dep_time`, `sched_dep_time` and `dep_delay`. How would you expect those three numbers to be related?

```
nycflights13::flights%>%
  dplyr::select(
    dep_time,
    sched_dep_time,
    dep_delay)
```

```
## # A tibble: 336,776 x 3
##   dep_time sched_dep_time dep_delay
##   <int>      <int>      <dbl>
## 1     517         515          2
## 2     533         529          4
## 3     542         540          2
## 4     544         545         -1
## 5     554         600         -6
## 6     554         558         -4
## 7     555         600         -5
## 8     557         600         -3
## 9     557         600         -3
## 10    558         600         -2
## # ... with 336,766 more rows
```

They are related in the following form:

```
nycflights13::flights%>%
  dplyr::select(
    dep_time,
    sched_dep_time,
    dep_delay)%>%
  dplyr::mutate(
    dep_time_min = (dep_time%%100)*60+(dep_time%100),
    sched_dep_time_min = (sched_dep_time%%100)*60+(sched_dep_time%100),
    calc=dep_time_min-sched_dep_time_min)
```

```
## # A tibble: 336,776 x 6
##   dep_time sched_dep_time dep_delay dep_time_min sched_dep_time_min calc
##   <int>      <int>      <dbl>      <dbl>          <dbl>      <dbl>
## 1     517         515          2          317            315          2
## 2     533         529          4          333            329          4
## 3     542         540          2          342            340          2
## 4     544         545         -1          344            345         -1
## 5     554         600         -6          354            360         -6
## 6     554         558         -4          354            358         -4
## 7     555         600         -5          355            360         -5
## 8     557         600         -3          357            360         -3
## 9     557         600         -3          357            360         -3
## 10    558         600         -2          358            360         -2
## # ... with 336,766 more rows
```

### 4

Find the 10 most delayed flights using a ranking function. How do you want to handle ties? Carefully read the documentation for `min_rank()`

```
nycflights13::flights%>%
  select(
    flight,
    dep_delay)->delays
```

I don't know what exactly is happening here with the `min_rank` function.

5

What does `1:3+1:10` return? Why?

```
1:3+1:10
```

```
## Warning in 1:3 + 1:10: longer object length is not a multiple of shorter object
## length
```

```
## [1] 2 4 6 5 7 9 8 10 12 11
```

Since the vectors are of different length the smallest vector is repeated in the operation.

6

What trigonometric functions does R provide?

The basic three functions are: `* cos(x)`

`* sin(x)` `* tan(x)`

These three following functions calculate the arc and the two-argument arc-tan:

- `acos(x)`
- `asin(x)`
- `atan(x)`
- `atan2(x,y)`

The last three functions calculate the `function(pi*x)`.

- `cospi(x)`
- `sinpi(x)`
- `tanpi(x)`

## Grouped summaries with `summarise()`


```
nycflights13::flights%>%
  dplyr::group_by(year, month, day)%>%
  dplyr::summarise(
    delay = mean(dep_delay, na.rm=TRUE),
    .groups = "drop")->gp
```

```
nycflights13::flights%>%
  dplyr::group_by(dest)%>%
  dplyr::summarise(
    count = n(),
    dist = mean(distance, na.rm=TRUE),
    delay = mean(arr_delay, na.rm=TRUE)
  )->dest #just to visualize the dataframe
```

```
nycflights13::flights%>%
  dplyr::group_by(dest)%>%
```

```
dplyr::summarise(
  count = n(),
  dist = mean(distance, na.rm=TRUE),
  delay = mean(arr_delay, na.rm=TRUE))%>%
dplyr::filter(count > 20, dest != "HNL")%>%
ggplot2::ggplot(aes(x=dist, y=delay))+
  geom_point()+
  geom_smooth()
```


```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
nycflights13::flights%>%
  dplyr::group_by(tailnum)%>%
  summarise(
    delay = mean(arr_delay, na.rm=TRUE)
  )%>%
  ggplot2::ggplot(aes(x=delay))+
  geom_freqpoly(binwidth=10)
```

```
## Warning: Removed 7 rows containing non-finite values (stat_bin).
```






Number of flights per average delay:

```
nycflights13::flights%>%  
  dplyr::group_by(tailnum)%>%  
  summarise(  
    delay = mean(arr_delay, na.rm=TRUE),  
    n=n())%>%  
  dplyr::filter(n > 25)%>%  
  ggplot2::ggplot(aes(x=n, y=delay))+  
  geom_point(alpha = 1/10)
```

## Warning: Removed 1 rows containing missing values (geom\_point).



Counts and proportions of logical values: `sum(x > 10)`, `mean(y == 0)`. When used with numeric functions, TRUE is converted to 1 and FALSE to 0. This makes `sum()` and `mean()` very useful: `sum(x)` gives the number of TRUEs in `x`, and `mean(x)` gives the proportion.

## Exercises 5.6.7

1

Brainstorm at least 5 different ways to assess the typical delay characteristics of a group of flights:

- A flight is 15 minutes early 50% of the time, and 15 minutes late 50% of the time.
- A flight is always 10 minutes late.
- A flight is 30 minutes early 50% of the time, and 30 minutes late 50% of the time.
- 99% of the time a flight is on time. 1% of the time it's 2 hours late.

Which is more important: arrival delay or departure delay?

2

Come up with another approach that will give you the same output as `not_cancelled%>%count(dest)` and `not_cancelled%>%count(tailnum, wt=distance)` (without using `count()`).

Lets check the output from the book:

```
nycflights13::flights%>%
  dplyr::filter(!is.na(dep_delay), !is.na(arr_delay))>%not_cancelled
```

```
not_cancelled %>%
  count(dest)->nc

not_cancelled %>%
  count(tailnum, wt = distance)->wtd
```

The other way to achieve the same results is grouping by and then summarizing:

```
not_cancelled%>%
  dplyr::group_by(dest)%>%
  dplyr::summarise(
    n=n()
  )->my_nc

not_cancelled%>%
  dplyr::group_by(tailnum)%>%
  summarise(
    totalmiles = sum(distance)
  )->my_wtd
```

### 3

Our definition of cancelled flights (`is.na(dep_delay)|is.na(arr_dealy)`) is slightly sub optimal. Why? Which is the most important column?

This is suboptimal because the command is assessing two vectors. Instead I propose using just the variable Air Time. We can check the following code:

```
nycflights13::flights%>%
  dplyr::select(
    dep_delay,
    arr_delay,
    air_time,
    dplyr::everything())%>%
  dplyr::filter(is.na(air_time))->arrg
```

Filtering for just the non-available numbers in the air time variable, we can see that is basically the same result as getting the n.a. values from departure delay **or** n.a. values from arrival delay (I believe: `is.na(air_time)==is.na(dep_delay)|is.na(arr_dealy)`)


### 4

Look at the number of cancelled flights per day. Is there a pattern? Is the proportion of cancelled flights related to the average delay?

Even if is suboptimal, I will use the book's definition of cancelled flights:


```
nycflights13::flights%>%
  dplyr::filter(is.na(dep_delay)|is.na(arr_delay))%>%
  dplyr::group_by(day)%>%
  dplyr::summarise(
    avg_delay=mean(dep_delay, na.rm=TRUE),
    n=n())%>%
  ggplot2::ggplot(aes(x=n, y=avg_delay))+
  geom_point()+
  geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
nycflights13::flights%>%  
  dplyr::filter(is.na(dep_delay)|is.na(arr_delay))%>%  
  dplyr::group_by(day)%>%  
  dplyr::summarise(  
    n=n())%>%  
  ggplot2::ggplot(aes(x=day,y=n))+  
    geom_point()+  
    geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



It does not seem to be a clear pattern.

5

Which carrier has the worst delays? Challenge: can you disentangle the effects of bad airport vs. bad carriers? Why/why not? (Hint: think about `flights%>%group_by(carrier,dest)%>%summarise(n())` )

```
nycflights13::flights%>%
  dplyr::group_by(carrier)%>%
  summarise(
    avg_dep_delay=mean(dep_delay, na.rm=TRUE))%>%
  dplyr::arrange(desc(avg_dep_delay))
```

```
## # A tibble: 16 x 2
##   carrier avg_dep_delay
##   <chr>      <dbl>
## 1 F9         20.2
## 2 EV         20.0
## 3 YV         19.0
## 4 FL         18.7
## 5 WN         17.7
## 6 9E         16.7
## 7 B6         13.0
## 8 VX         12.9
## 9 OO         12.6
## 10 UA        12.1
## 11 MQ        10.6
## 12 DL         9.26
```

```
## 13 AA      8.59
## 14 AS      5.80
## 15 HA      4.90
## 16 US      3.78

nycflights13::flights%>%
  dplyr::group_by(carrier)%>%
  summarise(
    avg_arr_delay=mean(arr_delay, na.rm=TRUE))%>%
  dplyr::arrange(desc(avg_arr_delay))

## # A tibble: 16 x 2
##   carrier avg_arr_delay
##   <chr>      <dbl>
## 1 F9         21.9
## 2 FL         20.1
## 3 EV         15.8
## 4 YV         15.6
## 5 OO         11.9
## 6 MQ         10.8
## 7 WN          9.65
## 8 B6          9.46
## 9 9E          7.38
## 10 UA         3.56
## 11 US         2.13
## 12 VX         1.76
## 13 DL         1.64
## 14 AA         0.364
## 15 HA        -6.92
## 16 AS        -9.93
```

F9 (frontier) is the worst carrier for both arrival and departure delays.

Now, taking the hint from the book:

```
nycflights13::flights%>%
  dplyr::group_by(carrier, dest)%>%
  summarise(n=n())->destcar
```

```
## `summarise()` has grouped output by 'carrier'. You can override using the
## `.groups` argument.
```

If we filter the flights made by frontier, we can see that they operate only between La Guardia (LGA) and Denver (DEN). Denver is one of the busiest airports in the US, it is possible that all flights related to Denver have big delays.

## 6

What does the `sort` argument do to `count()`. When might you use it?

When `sort=TRUE`, according to the documentation, the largest groups will appear at the top. It is not clear for me when this would be most useful.

```
People<-tibble::tibble(
  number=c(1,2,3,4,5,6),
  group=c(1,2,1,2,1,2))
```

## Grouped Mutates (and filters)

### Exercises 5.7.1

1

Refer back to the lists of useful mutate and filtering functions. Describe how each operation changes when you combine it with grouping.

2

Which plane (`tailnum`) has the worst on-time record?

3

What time of day should you fly if you want to avoid delays as much as possible?

4

For each destination, compute the total minutes of delay. For each flight, compute the proportion of the total delay for its destination.

5

Delays are typically temporally correlated: even once the problem that caused the initial delay has been resolved, later flights are delayed to allow earlier flights to leave. Using `lag()`, explore how the delay of a flight is related to the delay of the immediately preceding flight.

6

Look at each destination. Can you find flights that are suspiciously fast? (i.e. flights that represent a potential data entry error). Compute the air time of a flight relative to the shortest flight to that destination. Which flights were most delayed in the air?

7

Find all destinations that are flown by at least two carriers. Use that information to rank the carriers.

8

For each plane, count the number of flights before the first delay of greater than 1 hour.